

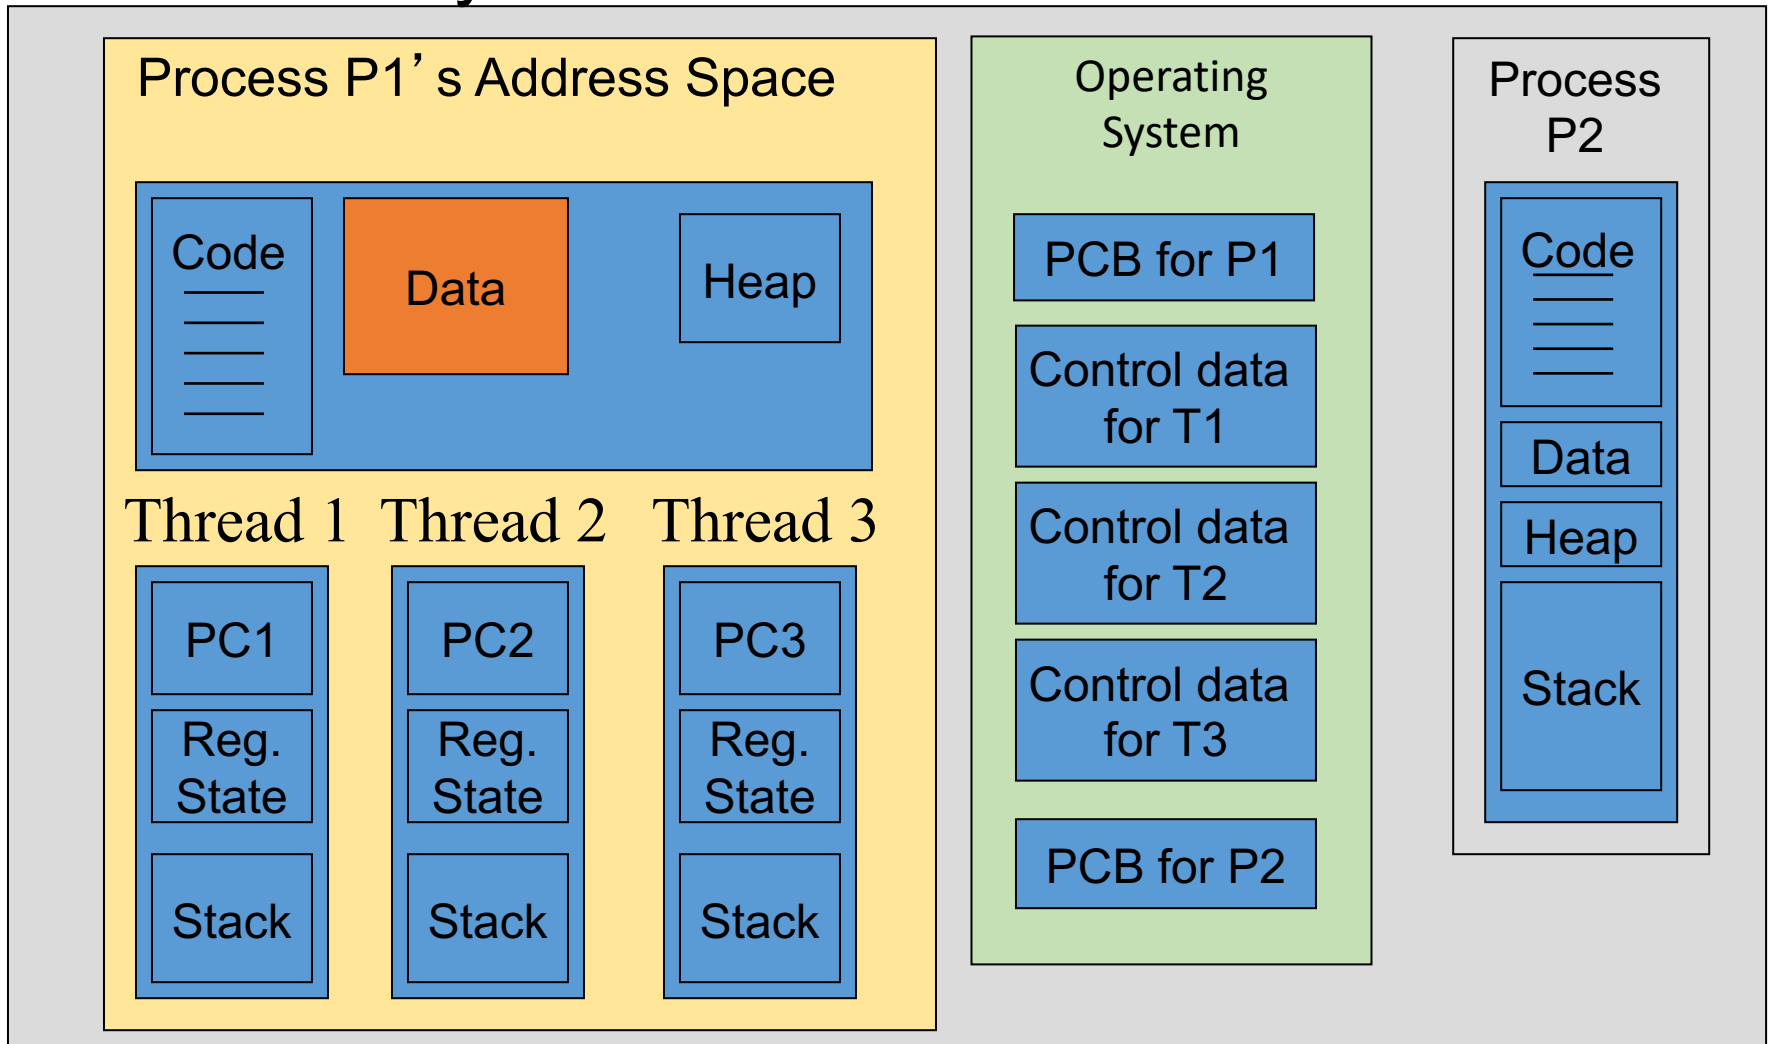


Lecture 7

Threads

Multiple Threads

Main Memory



Why do we want to use Threads

- Reduced context switch overhead vs multiple processes
 - E.g. In Solaris, context switching between processes is 5x slower than switching between threads
 - Don't have to save/restore context, including base and limit registers and other MMU registers, also TLB cache doesn't have to be flushed
- Shared resources => less memory consumption
 - Don't duplicate code, data or heap or have multiple PCBs as for multiple processes
 - Supports more threads – more scalable, e.g. Web server must handle thousands of connections



More reasons for

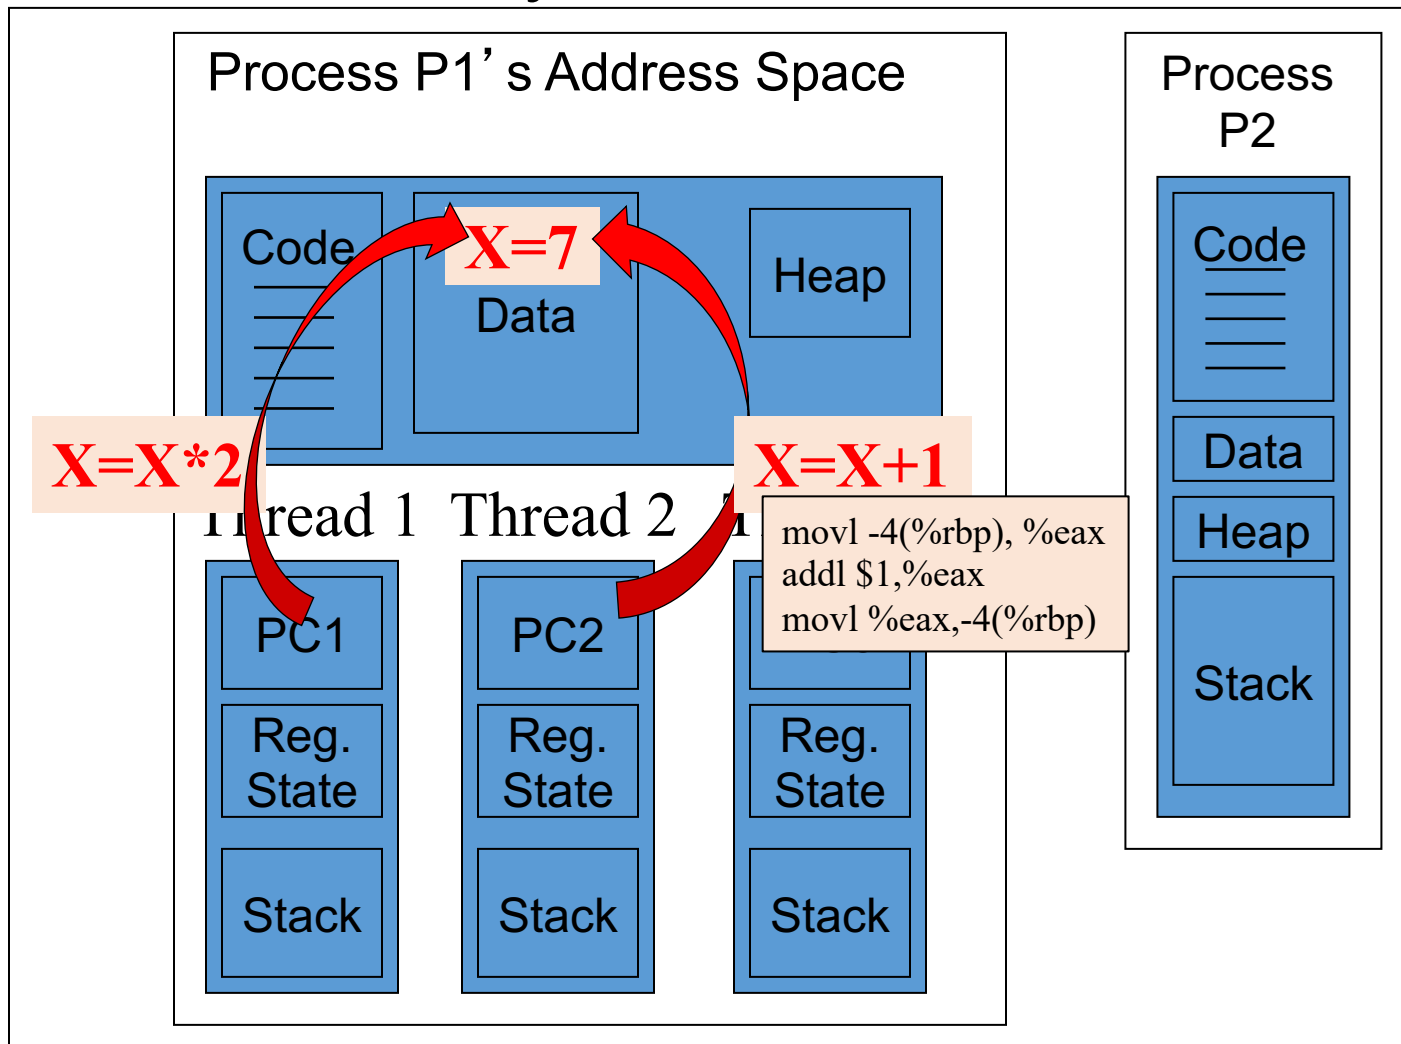
Why do we want to use Threads

- Inter-thread communication is easier and faster than inter-process communication
 - threads share the same memory space, so just read/write from/to the same memory location !!!
 - IPC via message passing uses system calls to send/receive a message, which is slow
 - IPC using shared memory may be comparable to inter-thread communication



Thread Safety

Main Memory



Suppose:

- Thread1 wants to multiply X by 2
- Thread2 wants to increment X
- *Could have a race condition (see Chapter 5)*

Thread-Safe Code

- A piece of code is **thread-safe** if it functions correctly during simultaneous or *concurrent* execution by multiple threads.
 - In particular, it must satisfy the need for multiple threads to access the same shared data, and the need for a shared piece of data to be accessed by only one thread at any given time.
- If two threads share and execute the same code, then unprotected use of shared
 - global variables is not thread safe
 - static variables is not thread safe
 - heap variables is not thread safe

We will learn how to write thread-safe code in Chapter 5

Processes vs. Threads

- Why are processes still used when threads bring so many advantages?
 1. Some tasks are sequential and not easily parallelizable, and hence are single-threaded by nature
 2. No fault isolation between threads
 - If a thread crashes, it can bring down other threads
 - If a process crashes, other processes continue to execute, because each process operates within its own address space, and so one crashing has limited effect on another
 - **Caveat:** a crashed process may fail to release synchronization locks, open files, etc., thus affecting other processes . But, the OS can use PCB's information to help cleanly recover from a crash and free up resources.

Processes vs. Threads (2)

- Why are processes still used when threads bring so many advantages? (cont.)
 3. Writing thread-safe/reentrant code is difficult. Processes can avoid this by having separate address spaces and separate copies of the data and heap



Threads vs. Processes

- Advantages of multithreading
 - Sharing between threads is easy
 - Faster creation
- Disadvantages of multithreading
 - Ensure threads-safety
 - Bug in one thread can bleed to other threads, since they share the same address space
 - Threads must compete for memory
- Considerations
 - Dealing with signals in threads is tricky
 - All threads must run the same program
 - Sharing of files, users, etc



Applications, Processes, and Threads

- An application can consist of multiple processes, each one dedicated to a specific task (UI, computation, communication, etc.)
- Each process can consists of multiple threads
- An application could thus consist of many processes and threads



Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS
- Three main thread libraries in use today:
 - POSIX Pthreads
 - Win32
 - Java



Thread Libraries

- Three main thread libraries in use today:

- **POSIX pthreads**

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library

- **Win32**

- Kernel-level library on Windows system

- **Java**

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS



The pthreads API

- ***Thread management:*** The first class of functions work directly on threads - creating, terminating, joining, etc.
- ***Semaphores:*** provide for create, destroy, wait, and post on semaphores.
- ***Mutexes:*** provide for creating, destroying, locking and unlocking mutexes.
- ***Condition variables:*** include functions to create, destroy, wait and signal based upon specified variable values.

Thread Creation

pthread_create (tid, attr, start_routine, arg)

- It returns the new thread ID via the *tid* argument.
- The *attr* parameter is used to set thread attributes, NULL for the default values.
- The *start_routine* is the C routine that the thread will execute once it is created.
- A single argument may be passed to *start_routine* via *arg*. It must be passed by reference as a pointer cast of type void.

Thread Termination and Join

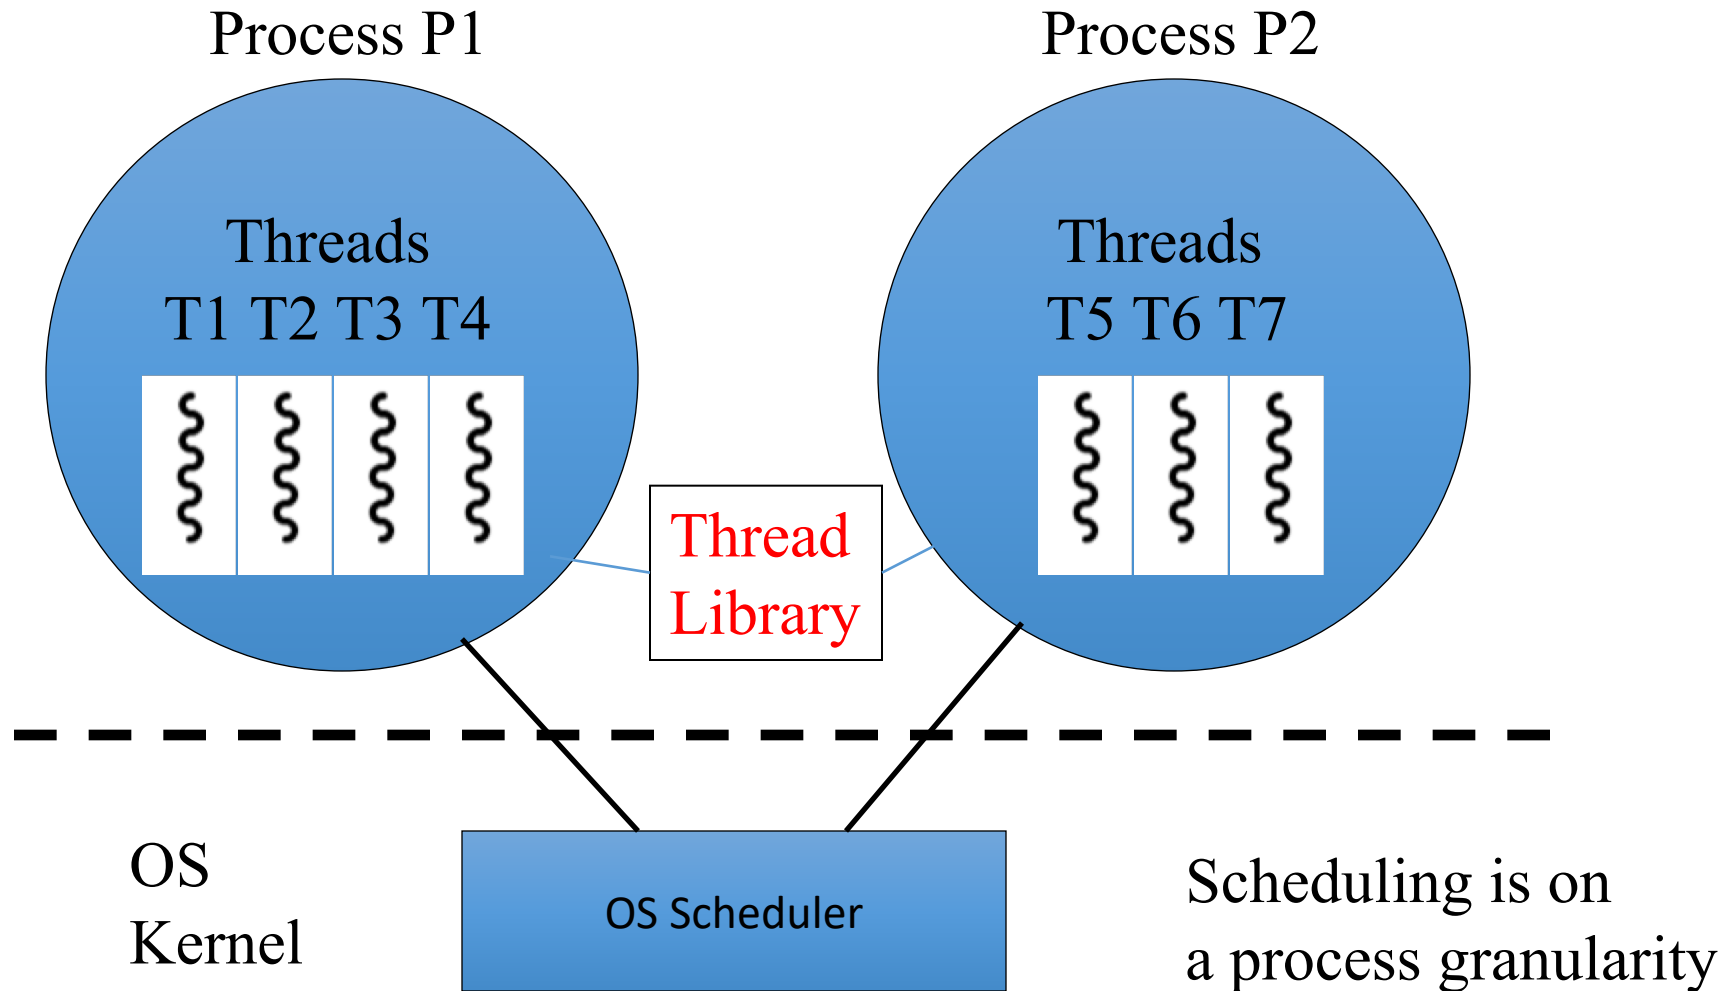
pthread_exit (value) ;

- This Function is used by a thread to terminate.
- The return value is passed as a pointer.

pthread_join (tid, value_ptr);

- The pthread_join() subroutine blocks the calling thread until the specified *threadid* thread terminates.
- Return 0 on success, and negative on failure. The returned value is a pointer returned by reference. If you do not care about the return value, you can pass NULL for the second argument.

User-Space Threads

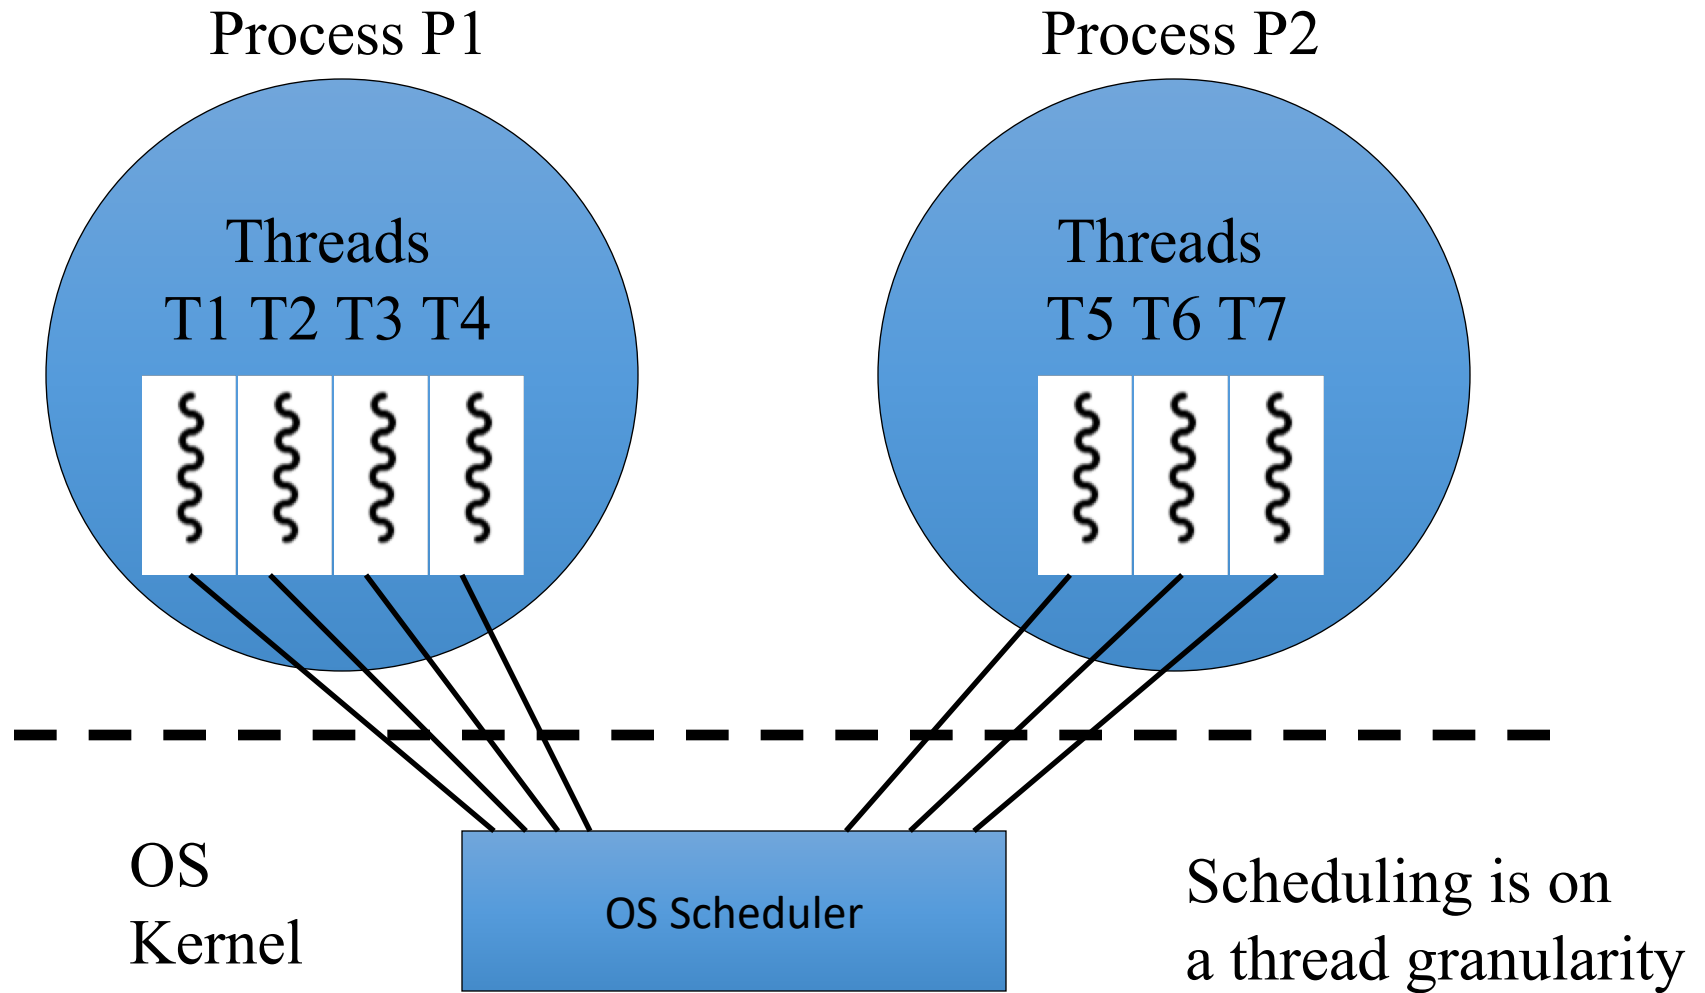


User-Space Threads

- *User space threads* are usually cooperatively multitasked, i.e. user threads within a process voluntarily give up the CPU to each other
 - threads will synchronize with each other via the user space threading package or library
 - Thread library: provides interface to create, delete threads in the same process
- OS is unaware of user-space threads – only sees user-space processes
 - If one user space thread blocks, the entire process blocks in a many-to-one scenario (see text)
- *pthread*s is a POSIX threading API
 - implementations of pthreads API differ underneath the API; could be user space threads; there is also pthreads support for kernel threads as well
- User space thread also called a *fiber*



Kernel Threads




Kernel Threads


- *Kernel threads* are supported by the OS
 - kernel sees threads and schedules at the granularity of threads
 - Most modern OSs like Linux, Mac OS X, Win XP support kernel threads
 - Mapping of user-level threads to kernel threads is usually one-to-one, e.g. Linux and Windows, but could be many-to-one, or many-to-many
 - Win32 thread library is a kernel-level thread library



Many-to-One Model



One-to-one Model



Many-to-Many Model

