



Lecture 10

Semaphores

Semaphores

- Dijkstra proposed more general solution to mutual exclusion
- Semaphore **S is an abstract data type** that is accessed only through two standard atomic operations
 - **wait() (also called P(), from Dutch word *proberen* “to test”)**
 - somewhat equivalent to a test-and-set
 - also involves *decrementing* the value of S
 - **signal() (V(), from Dutch word *verhogen* “to increment”)**
 - *increments* the value of S
- OS provides ways to create and manipulate semaphores atomically



Semaphores

```
typedef struct {  
    int value;  
    PID *list[ ];  
} semaphore;
```

Both wait() and signal()
operations are atomic

```
wait(semaphore *s) {  
    s→value--;  
    if (s→value < 0) {  
        add this process to s→list;  
        sleep ( );  
    }  
}
```

```
signal(semaphore *s) {  
    s→value++;  
    if (s→value <= 0) {  
        remove a process P from s→list;  
        wakeup (P);  
    }  
}
```



Binary Semaphore

```
semaphore S = 1; // initial value of semaphore is 1  
int counter;      // assume counter is set correctly somewhere in code
```

Process P1:

```
wait(S);  
    // execute critical section  
    counter++;  
signal(S);
```

Process P2:

```
wait(S);  
    // execute critical section  
    counter--;  
signal(S);
```

- Both processes atomically wait() and signal() the semaphore S, which enables mutual exclusion on critical section code, in this case protecting access to the shared variable *counter*
- ***This solves mutual exclusion but will also eliminate busy wait***

Problems with semaphores

shared R1, R2;

semaphore Q = 1; // binary semaphore as a mutex lock for R1

semaphore S = 1; // binary semaphore as a mutex lock for R2

Process P1:

wait(S); (1)

wait(Q); (3)

modify R1 and R2;

signal(S);

signal(Q);

Process P2:

wait(Q); (2)

wait(S); (4)

modify R1 and R2;

signal(Q);

signal(S);

Potential for deadlock

Deadlock

- In the previous example,
 - Each process will block on a semaphore
 - The signal() statements will never get executed, so there is no way to wake up the two processes
 - There is no rule about the order in which wait() and signal() operations may be invoked
- In general, with more processes sharing more semaphores, the potential for deadlock grows

Other problematic scenarios

- A programmer mistakenly follows a `wait()` with a second `wait()` instead of a `signal()`
- A programmer forgets and omits the `wait(mutex)` or `signal(mutex)`
- A programmer reverses the order of `wait()` and `signal()`

Another problem with synchronization

shared R1, R2;

semaphore S=1; // binary semaphore as a mutex lock for R1 & R2

Process P1:

wait(S)

modify R1 and R2;

Signal (S);

Process P2:

wait(S)

modify R1 and R2;

signal(S);

Process P3:

wait(S);

modify R1 and R2;

signal(S);

Potential for starvation

Starvation

- The possibility that a process would never get to run
- For example, in a multi-tasking system the resources could switch between two individual processes
- Depending on how the processes are scheduled, a third process may never get to run
- The third task is being starved of accessing the resource



Semaphore Solution for Mutual Exclusion

```
shared lock = 1; // initial value of semaphore 1
shared int count;
Shared data_type buffer [MAX];
```

Code for p₁

```
while(1) {
    produce (nextdata)
    while(count==MAX);
    wait(lock);
    buffer[count] = nextdata;
    count++;
    signal(lock);
}
```

Code for p₂

```
while (1) {
    while(count==0);
    wait(lock);
    data = buffer[count-1];
    count--;
    signal(lock);
    consume(data);
}
```



Semaphore Solution for Mutual Exclusion

```
shared lock = 1; // initial value of semaphore 1
shared int count;
Shared data_type buffer [MAX];
```

Code for p₁

```
while(1) {
    produce (nextdata)
    while(count==MAX);
    wait(lock);
    buffer[count] = nextdata;
    count++;
    signal(lock);
}
```

Code for p₂

```
while (1) {
    while(count==0);
    wait(lock);
    data = buffer[count-1];
    count--;
    signal(lock);
    consume(data);
}
```

- Busy waiting removed from the mutual exclusion when waiting on lock
- Does it solve all busy waiting issues?



pthread Synchronization

- Mutex locks
 - Used to protect critical sections
- Some implementations provide semaphores through POSIX SEM extension
 - Not part of pthread standard

```
#include <pthread.h>
pthread_mutex_t m; //declare a mutex object
pthread_mutex_init (&m, NULL); // initialize mutex object
```

```
//thread 1
pthread_mutex_lock (&m);
    //critical section code for th1
pthread_mutex_unlock (&m);
```

```
//thread 2
pthread_mutex_lock (&m);
    //critical section code for th2
pthread_mutex_unlock (&m);
```



pthread mutex

- pthread mutexes can have only one of two states: **lock** or **unlock**
- Important restriction
 - **Mutex ownership:**
Only the thread that locks a mutex can unlock that mutex
 - **Mutexes are strictly used for mutual exclusion while binary semaphores can also be used for synchronization between two threads or processes**



POSIX semaphores

```
#include <semaphore.h>
```

```
int sem_init(sem_t *sem, int pshared, unsigned int value);  
// pshared: 0 (among threads); 1 (among processes)
```

```
int sem_wait(sem_t *sem); //same as wait( )
```

```
int sem_post(sem_t *sem); //same as signal( )
```

```
sem_getvalue( ); // check the current value of the semaphore  
sem_close( ); // done with the semaphore then close
```

Producer-Consumer Problem


also known as

Bounded Buffer Problem

- We have already seen this problem with one producer and one consumer
- General problem: **multiple** producers and multiple consumers
- **Producers** puts new information in the buffer
- **Consumers** takes out information from the buffer




Prior Bounded-Buffer P/C Approach



- Producer places data into a buffer at the next available position
- Consumer takes information from the earliest item




Prior Bounded-Buffer P/C Approach



- Producer places data into a buffer at the next available position
- Consumer takes information from the earliest item

Prior Bounded-Buffer P/C Approach




```
while(1) {  
    produce (nextdata)  
    while(count==MAX);  
    Acquire(lock);  
    buffer[count] = nextdata;  
    count++;  
    Release(lock);  
}
```

```
while (1) {  
    while(count==0);  
    Acquire(lock);  
    data = buffer[count-1];  
    count--;  
    Release(lock);  
    consume(data);  
}
```



Prior Bounded-Buffer P/C Approach




```
while(1) {  
    produce (nextdata);  
    while(count==MAX);  
    Acquire(lock);  
    buffer[count] = nextdata;  
    count++;  
    Release(lock);  
}
```

```
while (1) {  
    while(count==0);  
    Acquire(lock);  
    data = buffer[count-1];  
    count--;  
    Release(lock);  
    consume(data);  
}
```

Busy-wait!



Prior Bounded-Buffer P/C Approach



```
while(1) {  
    produce (nextdata);  
    while(count==MAX);  
    Acquire(lock);  
    buffer[count] = nextdata;  
    count++;  
    Release(lock);  
}
```

while (TS(&lock));

Busy-wait!

```
while (1) {  
    while(count==0);  
    Acquire(lock);  
    data = buffer[count-1];  
    count--;  
    Release(lock);  
    consume(data);  
}
```

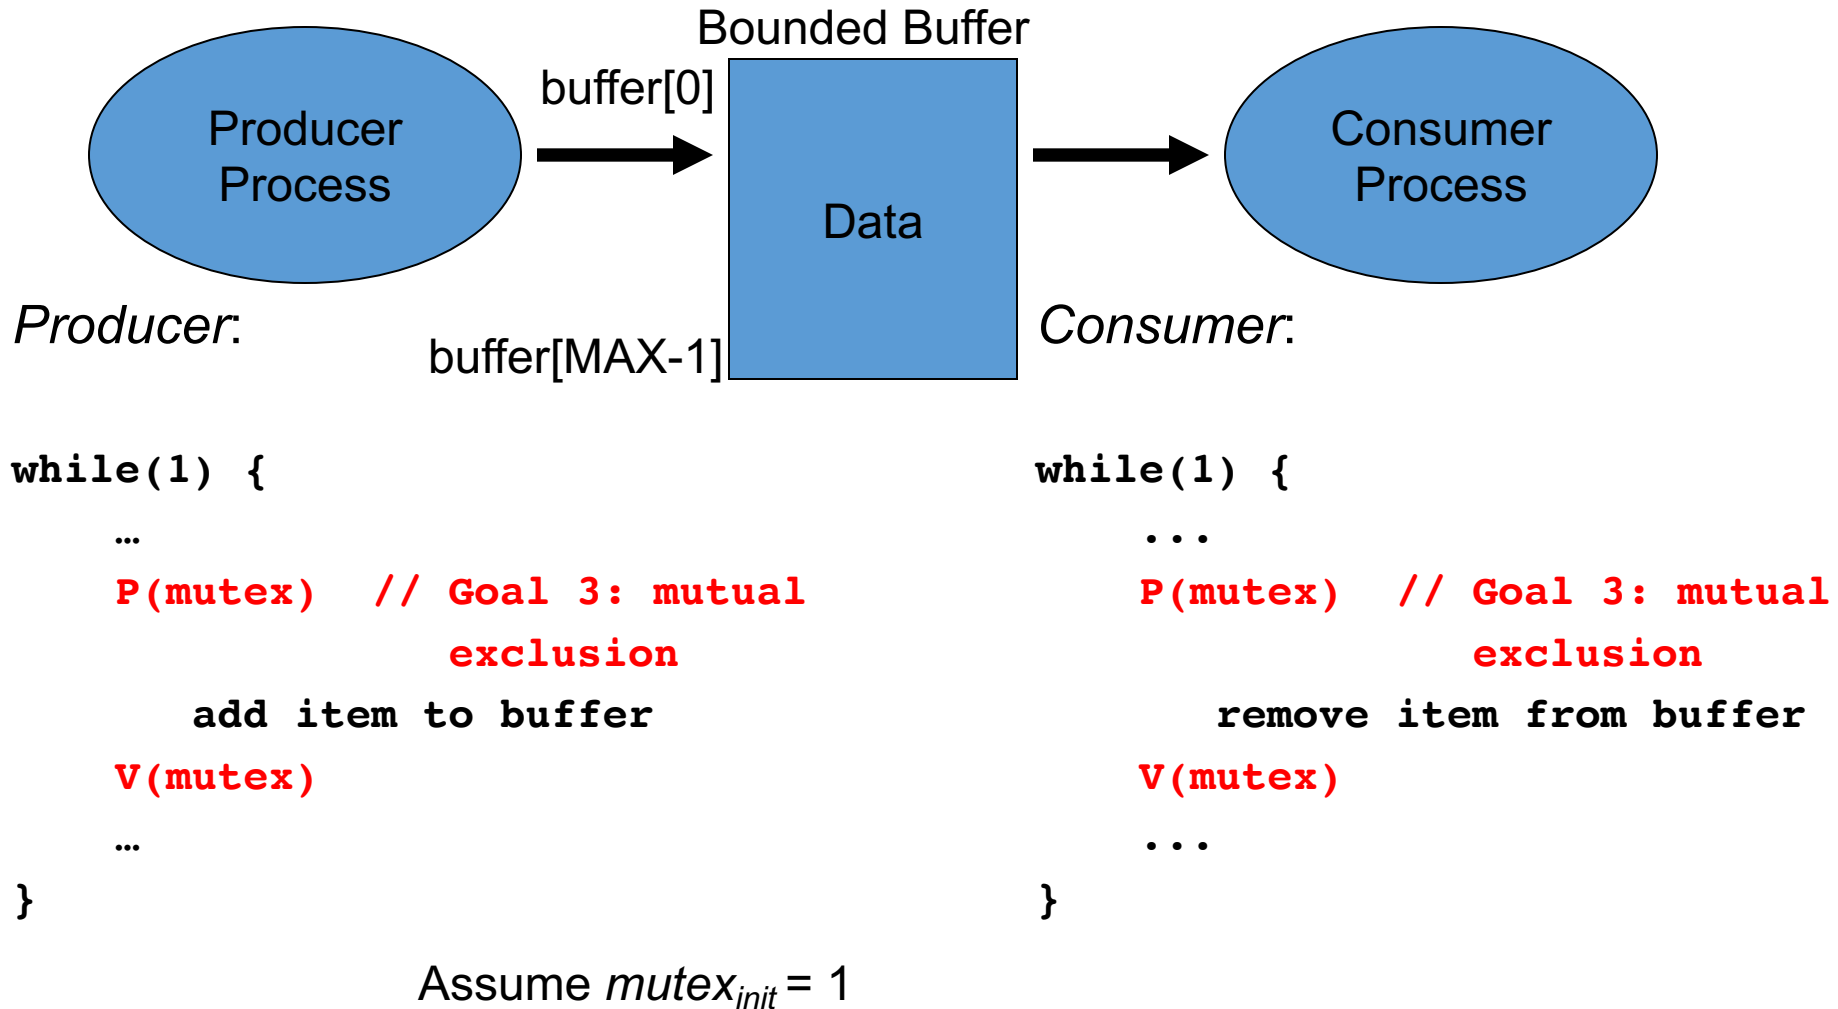


Bounded-Buffer Goals

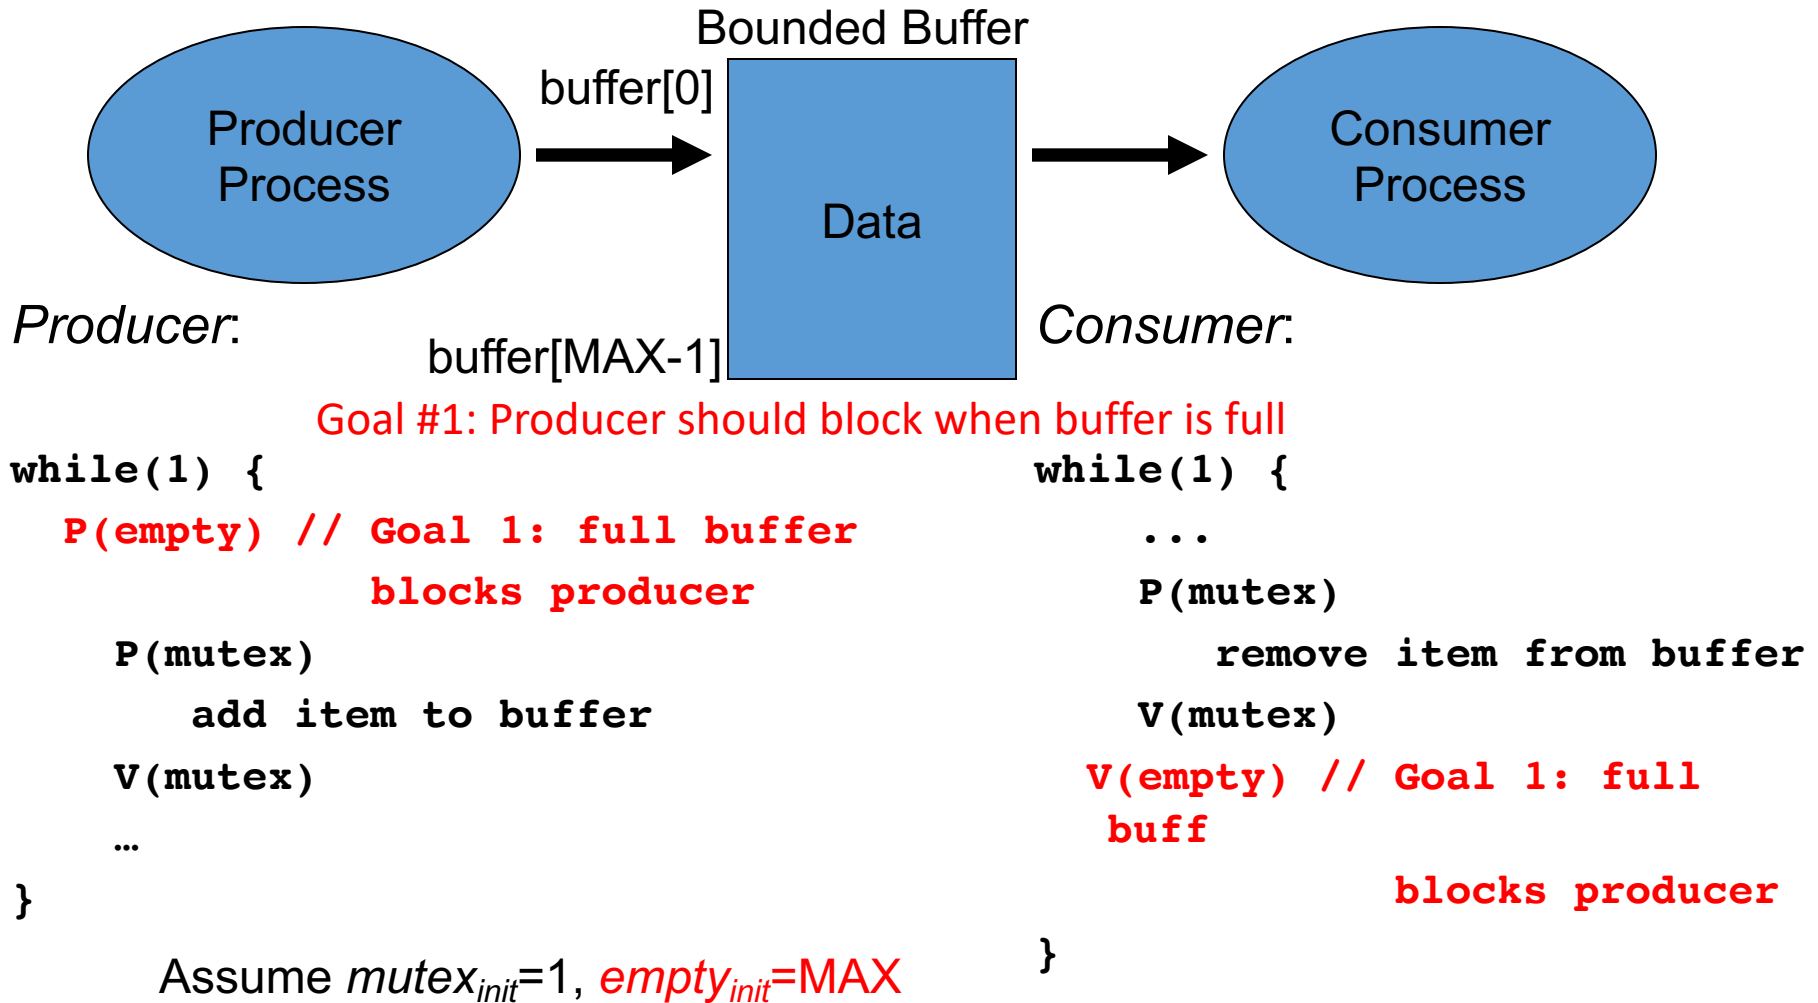
- In the prior approach, both the producer and consumer are ***busy-waiting*** using locks
- Instead, want both to sleep as necessary
 - Goal #1: Producer should block when buffer is full
 - Goal #2: Consumer should block when the buffer is empty
 - Goal #3: mutual exclusion when buffer is partially full
 - Producer and consumer should access the buffer in a synchronized mutually exclusive way



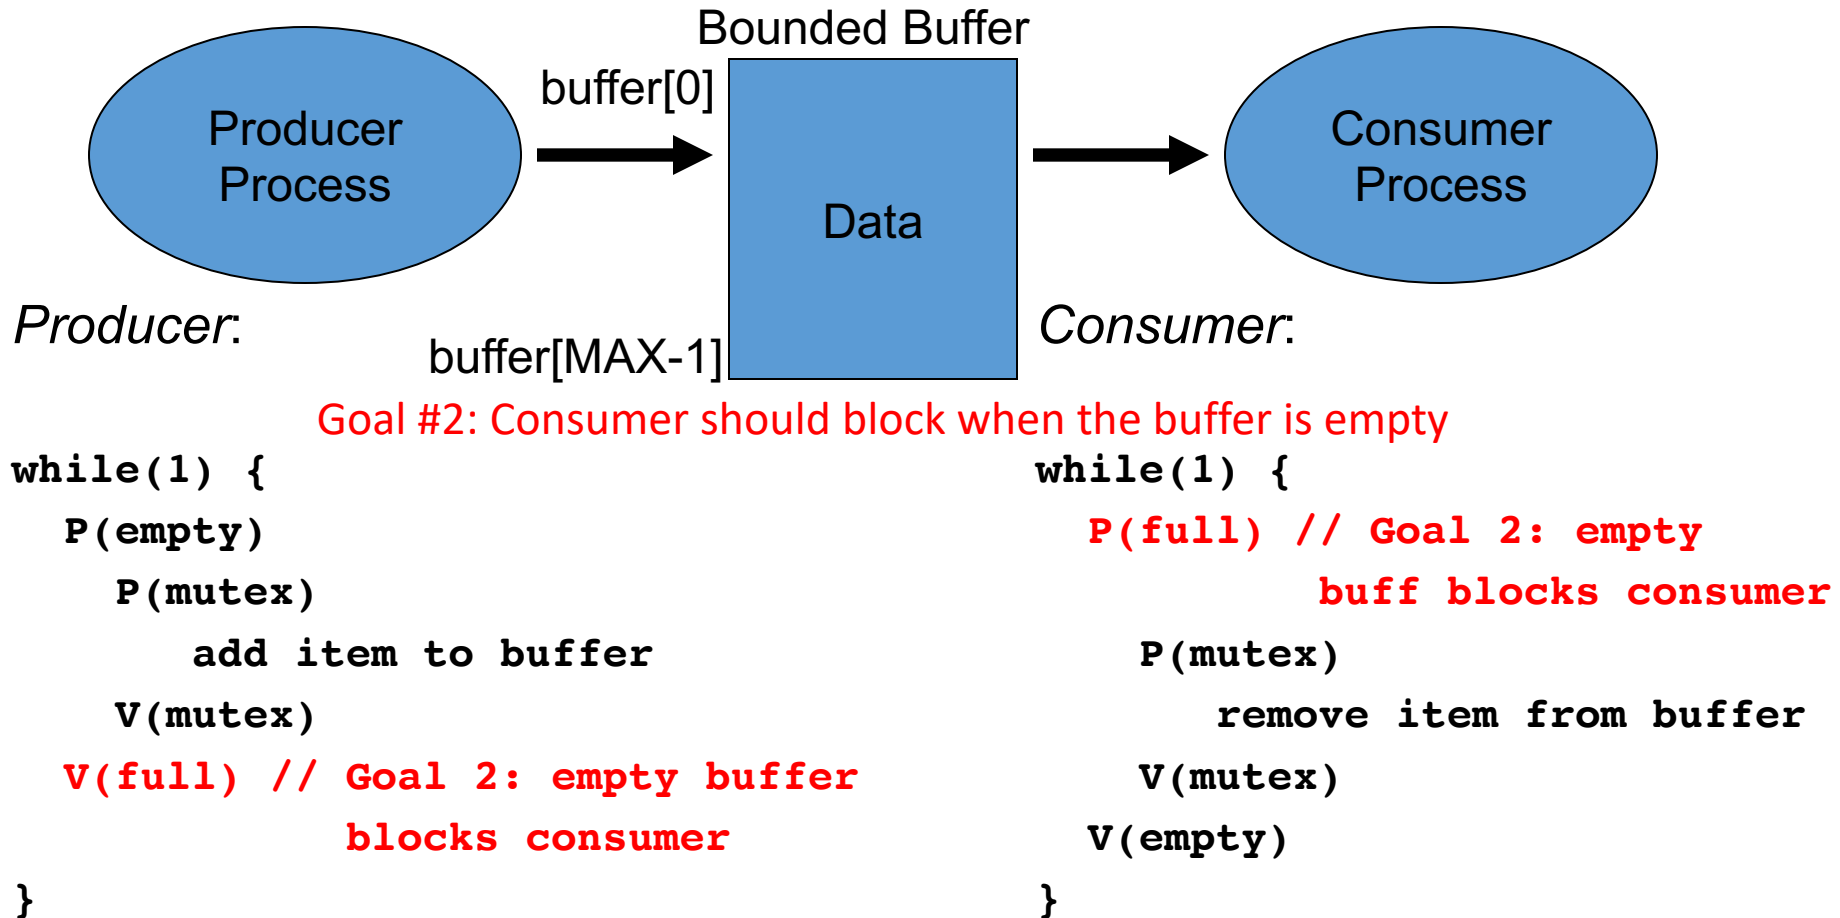
Bounded Buffer Solution



Bounded Buffer Solution (2)



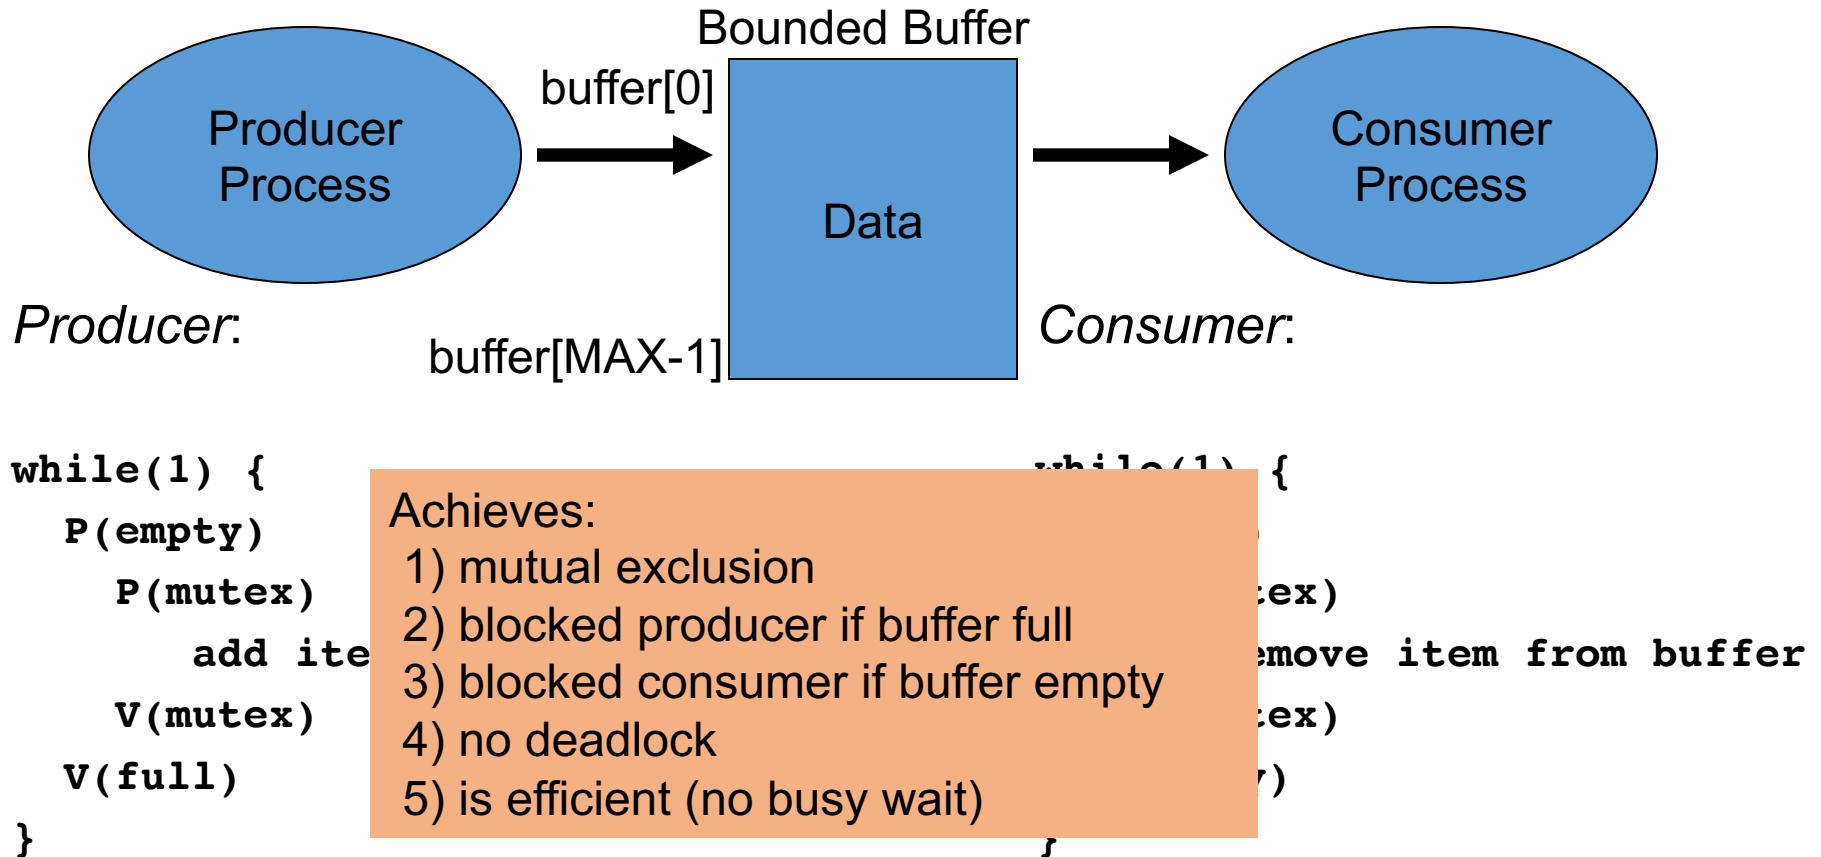
Bounded Buffer Solution (3)



Assume $mutex_{init}=1$, $empty_{init}=MAX$, $full_{init}=0$



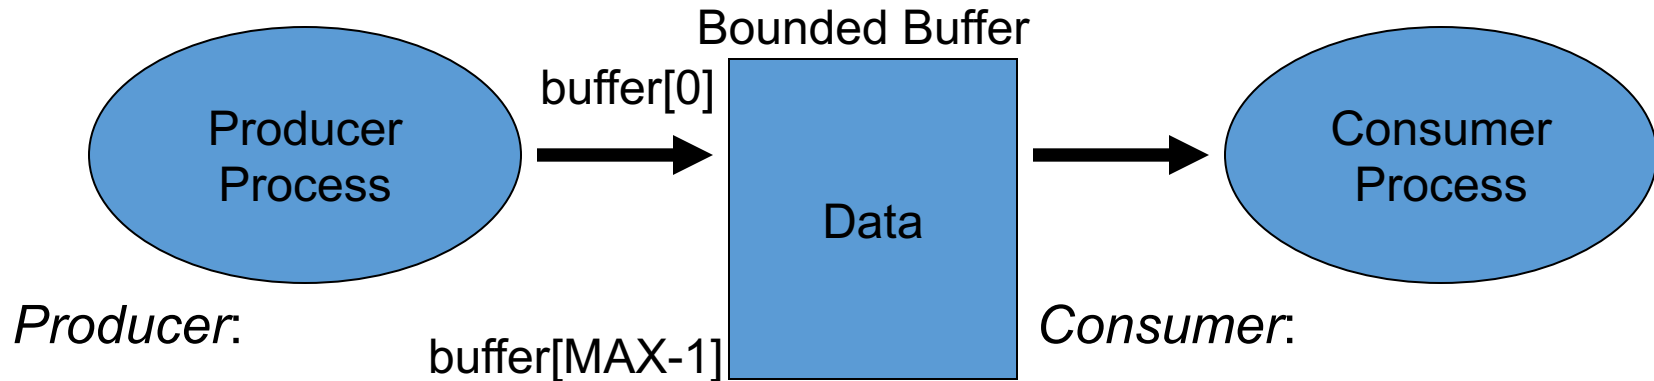
Bounded Buffer Solution (4)



Assume $mutex_{init}=1$, $empty_{init}=MAX$, $full_{init}=0$



Bounded Buffer Solution (5)



```
while(1) {  
    P(space_avail)  
    P(mutex)  
    add item to buffer  
    V(mutex)  
    V(items_avail)  
}  
    Assume  $mutex_{init}=1$ ,  
            $space\_avail_{init}=MAX$ ,  
            $items\_avail_{init}=0$ 
```

```
while(1) {  
    P(items_avail)  
    P(mutex)  
    remove item from buffer  
    V(mutex)  
    V(space_avail)  
}
```



Bounded-Buffer Design

- **Goal #1: Producer should block when buffer is full**
 - Use a counting semaphore called *empty* that is initialized to $empty_{init} = MAX$
 - Each time the producer adds an object to the buffer, this decrements the # of empty slots, until it hits 0 and the producer blocks
- **Goal #2: Consumer should block when the buffer is empty**
 - Define a counting semaphore *items_avail* that is initialized to $items_avail_{init} = 0$
 - *items_avail* tracks the # of full slots and is incremented by the producer
 - Each time the consumer removes a full slot, this decrements *items_avail*, until it hits 0, then the consumer blocks
- **Goal #3: Mutual exclusion when buffer is partially full**
 - Use a mutex semaphore to protect access to buffer manipulation, $mutex_{int} = 1$



Bounded Buffer Solution (6)

